

# ATS.ai

---

*AI-Powered Candidate Scoring Engine*

---

## Complete Project Report

Developed by **Krishna & Arya**

February 2026

 Live: <https://ats-ai-app.vercel.app>

 Source: <https://github.com/AryaJagtap/ATS.ai>

## 1. Problem Statement

---

Recruiters and HR teams manually screening resumes is slow, inconsistent, and doesn't scale. A company receiving 100+ resumes per job opening spends hours reading each one, often missing good candidates or applying inconsistent criteria.

### **Goal: Build an AI-powered ATS that:**

- ▶ Scores resumes against job descriptions using LLM analysis
- ▶ Handles 100+ resumes in minutes, not hours
- ▶ Provides actionable insights (best-fit role, missing skills, recommendations)
- ▶ Exports professional reports for HR teams

## 2. Starting Point — The Initial Codebase

---

The project began with a pre-existing codebase that had a FastAPI backend with basic resume scoring, a Next.js frontend with a basic UI, OpenAI GPT-4o as the primary LLM, pdfplumber for PDF extraction, CSV/XLSX file upload for candidate lists, and basic sequential processing (one resume at a time).

### **Initial Problems Identified**

Problem	Impact
Import errors (genai module)	Backend wouldn't start
Sequential processing	~5-6 seconds per resume, 100 resumes = 10+ minutes
No fallback when OpenAI fails	Rate limits (429 errors) caused complete failures
No single resume upload	Users had to create a spreadsheet even for 1 resume
PDF filename used as candidate name	"resume (1).pdf" shown instead of actual name
No multi-JD support	Only one job description at a time
Noisy terminal warnings	pdfminer font warnings cluttered logs

## 3. Development Phases — Step by Step

---

## Phase 1: Fixing Critical Errors

Problem: The backend had import errors related to the google-genai module. The import genai syntax was outdated.

Solution: Updated the import to use the correct module path: from google import genai. Fixed the initialization pattern to match the latest google-genai SDK.

What we gained: A working backend that could start without errors. What we lost: Nothing — purely a bug fix.

## Phase 2: Performance — Reducing Per-Resume Analysis Time

Each resume took 5-6 seconds to analyze. For 100 resumes, this meant 8-10 minutes of sequential processing.

### Time Breakdown Analysis

Step	Time	Bottleneck?
PDF download	~0.5s	No
Text extraction (pdfplumber)	~0.1-0.2s	No
Keyword scoring (TF-IDF)	~0.1s	No
LLM API call	~3-5s	YES — the bottleneck
Total	~5-6s	

### Optimizations Applied

- ▶ Switched LLM model (gpt-4o → gpt-4o-mini): Slightly less "creative" responses, but for structured JSON scoring output, quality is identical. 2-3x faster response times, significantly cheaper API costs. Clear win — no quality loss for our use case.
- ▶ Parallel keyword + LLM scoring: Used ThreadPoolExecutor to run keyword\_score and llm\_score simultaneously instead of sequentially. Saved ~0.1-0.2s per resume.
- ▶ Reduced LLM token overhead: Added max\_tokens=500 to limit response length. Set temperature=0 for deterministic results. Compacted the prompt to reduce input token count.
- ▶ Reusable LLM client singletons: Created single OpenAI/Gemini client instances instead of re-creating them per request. Eliminated connection setup overhead.
- ▶ Pre-imported heavy modules: Moved sklearn and regex imports to module level instead of inside functions. Saved ~0.1s on first call.

**Result: Per-resume time dropped from 5-6s → 2-3s ✓**

## Phase 3: Single Resume Upload

Problem: Users had to create a CSV/XLSX file even to analyze a single resume. Tedious for small tasks.

Solution: Added a toggle in the Upload Candidates section: "Single Resume" (drag & drop a single PDF/DOCX) and "Multiple Resumes" (upload CSV/XLSX or select multiple PDFs).

Files modified: FileUpload.js (new toggle UI), page.js (state management), main.py (new /api/analyze-direct endpoint).

What we gained: Much better UX for quick single-resume checks.

## Phase 4: Candidate Name Extraction from Resume

Problem: When a PDF was named generically (e.g., "resume (1).pdf"), the system showed the filename as the candidate name.

Solution: The LLM was already extracting the candidate name from resume content. Updated the logic to use the LLM-extracted name when the filename didn't appear to be a real name.

What we gained: Accurate candidate names from resume content.

## Phase 5: Multi-JD (Job Description) Matching

Problem: Companies often have multiple open positions. Users had to run separate analyses for each job description.

Solution: Enabled uploading multiple JD files. The system scores each candidate against ALL job descriptions and selects the best match.

### Implementation

- ▶ JobDescription.js — Updated to accept multiple file uploads
- ▶ main.py — Modified /api/analyze to iterate over multiple JDs per candidate
- ▶ ResultsTable.js — Added "Matched Role" column showing which JD was the best fit
- ▶ scorer.py — Score function runs against each JD, returns the highest score

Architecture decision: Score against ALL JDs and pick the best match automatically. More useful for HR teams, less manual work.

## Phase 6: Project Cleanup

The codebase had accumulated legacy files, unused dependencies, and outdated configs. Removed 14 legacy files and folders, updated requirements.txt with pinned versions, cleaned up .gitignore.

What we gained: Cleaner codebase, smaller deployment size.

## Phase 7: Batch Processing Optimization

Problem: 120 resumes took 10+ minutes even with per-resume optimizations.

### Evolution of Batch Size

Stage	Batch Size	120 Resumes Time	Issues
Initial	1 (sequential)	~15+ min	Way too slow
v1	5	~12 min	Still slow
v2	8	~10 min	Some rate limits
v3	12	~8 min	Occasional 429s
v4 (Final)	15	~6 min ✓	Retry logic handles 429s

Key decision: How high can we push batch concurrency? Constraint: OpenAI rate limits (RPM = Requests Per Minute). Tier 1 accounts: ~500 RPM for gpt-4o-mini. Solution: Batch size 15 + retry logic with exponential backoff. If a 429 hits, the system waits and retries (up to 2 times) instead of failing.

### Retry Logic Flow

- ▶ Attempt 1: Call OpenAI → If 429/500 error: wait 1 second
- ▶ Attempt 2: Retry OpenAI → If still fails: wait 2 seconds
- ▶ Attempt 3: Retry OpenAI → If still fails: fall back to Gemini (same retry pattern)
- ▶ If Gemini also fails: fall back to keyword-only scoring

**Result: 120 resumes in ~5 min 51 sec ✓ with 0 failures (down from 6 failures).**

## Phase 8: Error Handling & Terminal Cleanup

Problem: Terminal was flooded with pdfminer font warnings. Also, 6 out of 120 resumes were failing due to rate limits.

### Solutions Applied

- ▶ pdfminer warnings: Added logging suppression to hide non-critical font warnings
- ▶ Rate limit failures: Implemented retry logic with exponential backoff (Phase 7)
- ▶ TXT file support: Added .txt as a valid file type for job descriptions

## Phase 9: Time Elapsed Counter

Problem: Users had no visibility into how long the analysis was taking.

Solution: Implemented a dual-display timer:

- ▶ During processing: Live ticking "⌚ Time elapsed: 2m 15s" below the progress bar
- ▶ After completion: Static "⌚ Completed in 5m 51s" badge above the metrics cards

Bug found and fixed: Both timers showed simultaneously during processing. Fixed by gating the "Completed in" badge with a processing state check.

## Phase 10: PDF Parser Upgrade — PyMuPDF

### Parser Comparison

Parser	Speed	Quality	Memory	Our Choice
pdfplumber (was primary)	Moderate	Excellent	High	✗ Replaced
PyPDF2 (was fallback)	Slow	Good	Low	✓ Kept as fallback
PyMuPDF / fitz	~10x faster	Excellent	Moderate	✓ New primary
langextract (considered)	Unknown	Unknown	Unknown	✗ Not well-known

The user asked about langextract as an alternative. After analysis: (1) PDF extraction is NOT the bottleneck (0.1-0.2s vs 2-3s for LLM), (2) langextract is not a widely-used, well-documented package, (3) PyMuPDF is the industry-standard fastest PDF parser.

**Result: Switched to PyMuPDF as primary, PyPDF2 as fallback. Saved ~0.1s per resume. Removed pdfplumber dependency.**

## Phase 11: Photo Link Extraction from Spreadsheet

Problem: The input spreadsheet had a "Photograph" column with Google Drive photo URLs, but the output showed "Not Found" for all photos.

Root cause: The system was asking the LLM to find photo links inside the resume text. Resumes don't contain photo URL links — those links were in the spreadsheet itself.

Solution: Auto-detect columns named "Photo", "Photograph", "Image", or "Picture" in the input spreadsheet. Pass the URL through to the result, overriding the LLM's "Not Found".

Performance impact: Zero — just reading an existing column value.

## Phase 12: Favicon & Branding

Problem: The browser tab showed a generic emoji (🌐) as the favicon, and the project name was too long.

### Solutions

- ▶ Generated a custom AI/recruitment-themed logo (purple-to-magenta gradient)
- ▶ Placed as src/app/icon.png (Next.js App Router auto-serves this)
- ▶ Deleted old favicon.ico that was overriding custom icon

### Name Selection Process

Category	Names Considered
Short & Punchy	HireIQ, TalentLens, ScoreHire, ResumeAI, HireFlow
Techy / Modern	ATS.ai, Recruitr, MatchIQ, HireNex, TalentQ
Professional	ProHire, HireMetrics, SmartATS, CandidateIQ

Selected: ATS.ai — clean, memorable, domain-style, instantly communicates purpose.

## Phase 13: Full Rebranding to ATS.ai

File	Old Value	New Value
layout.js (tab)	AI Recruitment & ATS Platform	ATS.ai
Header.js (title)	AI Recruitment & ATS Platform	ATS.ai
Header.js (subtitle)	Fast, Rate-Limit Safe...	AI-Powered Candidate Scoring Engine
page.js (footer)	AI Recruitment & ATS Platform v2.0	ATS.ai
main.py (FastAPI)	AI Recruitment & ATS Platform	ATS.ai
README.md	AI Recruitment & ATS Platform v2.0	ATS.ai — AI-Powered Recruitment Platform

## Phase 14: Production Deployment

Component	Platform	Plan	URL
Frontend	Vercel	Free (Hobby)	ats-ai-app.vercel.app
Backend	Render	Free (Docker)	ats-ai-backend-kx9z.onrender.com

Component	Platform	Plan	URL
Source Code	GitHub	Public	github.com/AryaJagtap/ATS.ai

## Deployment Steps Completed

- ▶ Added MIT LICENSE file
- ▶ Initialized Git repository
- ▶ Pushed to GitHub (with Personal Access Token authentication)
- ▶ Created Render web service (Docker runtime, Singapore region)
- ▶ Set ALLOWED\_ORIGINS environment variable for CORS
- ▶ Created Vercel project (Root Directory: frontend)
- ▶ Set NEXT\_PUBLIC\_API\_URL pointing to Render backend
- ▶ Updated CORS to allow the real Vercel URL
- ▶ Verified end-to-end functionality

## 4. Technology Stack — Final

Layer	Technology	Version
Frontend Framework	Next.js	15.x
UI Library	React	19.x
Styling	CSS Variables (custom design system)	—
Backend Framework	FastAPI	0.131.0
Language	Python	3.10+
Primary LLM	OpenAI GPT-4o-mini	—
Fallback LLM	Google Gemini 2.5 Flash	—
Offline Scoring	scikit-learn TF-IDF + cosine similarity	1.8.0
PDF Parser (primary)	PyMuPDF (fitz)	1.27.1
PDF Parser (fallback)	PyPDF2	3.0.1
DOCX Parser	python-docx	1.2.0
Excel Export	openpyxl	3.1.5
Data Processing	pandas	3.0.1
File Download	gdown + requests	5.2.1 / 2.32.5
Frontend Hosting	Vercel	Free

Layer	Technology	Version
Backend Hosting	Render (Docker)	Free

## 5. Performance Journey

Metric	Start	Phase 2	Phase 7	Final
Per-resume time	5-6s	2-3s	2-3s	2-3s
120 resumes total	15+ min	~12 min	~8 min	5m 51s
Batch size	1	5	12	15
Failure rate	High	Moderate	6/120	0/120
LLM model	gpt-4o	gpt-4o-mini	gpt-4o-mini	gpt-4o-mini
PDF parser	pdfplumber	pdfplumber	pdfplumber	PyMuPDF

## 6. Key Lessons Learned

### What TO Use

Decision	Why
gpt-4o-mini over gpt-4o	2-3x faster, 10x cheaper, identical quality for structured JSON
PyMuPDF over pdfplumber	~10x faster for text extraction, smaller dependency
Batch concurrency (15)	Maximizes throughput without overwhelming API rate limits
Retry with exponential backoff	Handles transient API failures — zero resume failures
SSE (Server-Sent Events)	Real-time progress without WebSocket complexity
ThreadPoolExecutor	Simple, effective parallelism for I/O-bound LLM calls
CSS Variables	Clean dark/light theme toggle without library overhead

### What NOT to Use

Decision	Why
gpt-4o for scoring	Overkill — slower and expensive, no quality gain for JSON output

Decision	Why
pdfplumber as primary	Slower than PyMuPDF, heavier dependency (pulls in pdfminer)
langextract	Not well-known, no clear advantage, PDF parsing isn't the bottleneck
Sequential processing	Unacceptable for 100+ resumes — must use batch concurrency
WebSockets for progress	SSE is simpler, sufficient for one-directional progress updates
Hardcoded API keys	Security risk — use env vars or user-provided keys

## Key Tradeoffs Made

Tradeoff	Chose	Over	Reason
Speed vs accuracy	gpt-4o-mini (faster)	gpt-4o (slower)	No measurable accuracy loss for scoring
Batch size vs rate limits	15 + retries	Conservative batch of 5	Retries handle occasional 429s
API key model	User-provided keys	Developer pays for all	Sustainable — no cost to developer
Multi-JD scoring	Auto score all JDs	User selects per candidate	More automation, better UX
Photo extraction	Read from spreadsheet	Extract from resume via LLM	Resumes don't contain photo URLs

## 7. Architecture Flow

---

User uploads CSV/XLSX + Job Description(s)

### Frontend (Next.js / Vercel)

- ▶ File upload (drag & drop)
- ▶ JD text paste / file upload (multi-file)
- ▶ API key input via Settings
- ▶ SSE listener for real-time progress
- ▶ Timer (elapsed / completed)

↓ POST /api/analyze (multipart form data)

### Backend (FastAPI / Render)

- ▶ Parse CSV/XLSX → extract candidates (auto-detect Name, URL, Photo columns)
- ▶ For each batch of 15 candidates (ThreadPoolExecutor):
  - ▶ a. Download resume (gdown / requests)
  - ▶ b. Extract text (PyMuPDF primary → PyPDF2 fallback)
  - ▶ c. Score in parallel: Keyword + TF-IDF scoring (weight: 0.3) and LLM cascade: GPT-4o-mini → Gemini Flash → keyword-only (weight: 0.7)
  - ▶ d. If multi-JD: select best match across all JDs
  - ▶ e. Inject photo URL from spreadsheet
- ▶ Stream results via SSE (Server-Sent Events)
- ▶ Export to styled Excel report (openpyxl)

## 8. Final Feature List

#	Feature	Status
1	Multi-LLM cascade (GPT → Gemini → Keyword)	✓
2	ATS scoring (0-100, weighted blend)	✓
3	Real-time SSE progress streaming	✓
4	Dark / Light theme toggle	✓
5	Single resume upload	✓
6	Multiple resume upload (CSV/XLSX/PDFs)	✓
7	Multi-JD matching	✓
8	Photo link extraction from spreadsheet	✓
9	Candidate name extraction from resume	✓
10	Elapsed time counter	✓
11	Excel export (color-coded)	✓
12	API key config via UI	✓
13	Retry logic with exponential backoff	✓
14	Batch concurrency (15 parallel)	✓
15	Responsive design (desktop + mobile)	✓

#	Feature	Status
16	Custom favicon and branding	✓
17	Production deployment (Vercel + Render)	✓
18	MIT License	✓

## 9. Files Modified Summary

File	Changes Made
backend/main.py	Multi-JD endpoints, batch size 15, photo column detection, SSE streaming, branding
backend/utils/scorer.py	gpt-4o-mini, parallel scoring, retry logic, client singletons, compact prompt
backend/utils/extractor.py	PyMuPDF primary, PyPDF2 fallback, TXT support, warning suppression
backend/requirements.txt	Pinned versions, replaced pdfplumber with PyMuPDF
frontend/src/app/page.js	Timer, footer branding
frontend/src/app/layout.js	ATS.ai title, meta description, favicon
frontend/src/components/Header.js	ATS.ai branding
frontend/src/components/FileUpload.js	Single/Multiple resume toggle
frontend/src/components/JobDescription.js	Multi-file JD upload
frontend/src/components/ResultsTable.js	Matched Role column
README.md	Full rewrite with features, benchmarks, deployment guide
LICENSE	MIT License added

## 10. Conclusion

The ATS.ai project evolved from a basic resume scoring tool with critical errors into a production-grade, high-performance AI recruitment platform. Through systematic optimization — switching LLM models, implementing batch concurrency, adding retry logic, and upgrading PDF parsers — we achieved a 63% reduction in processing time (from 15+ minutes to ~6 minutes for 120 resumes) while simultaneously improving reliability from a ~5% failure rate to zero failures.

The platform is now deployed, branded, and ready for real-world use. Any user can bring their own API keys and start scoring resumes immediately.

---

— Project by **Krishna & Arya** —